

# Programming for persistent memory



## INTEL® PMDK WORKSHOP 英特尔® PMDK 研讨会

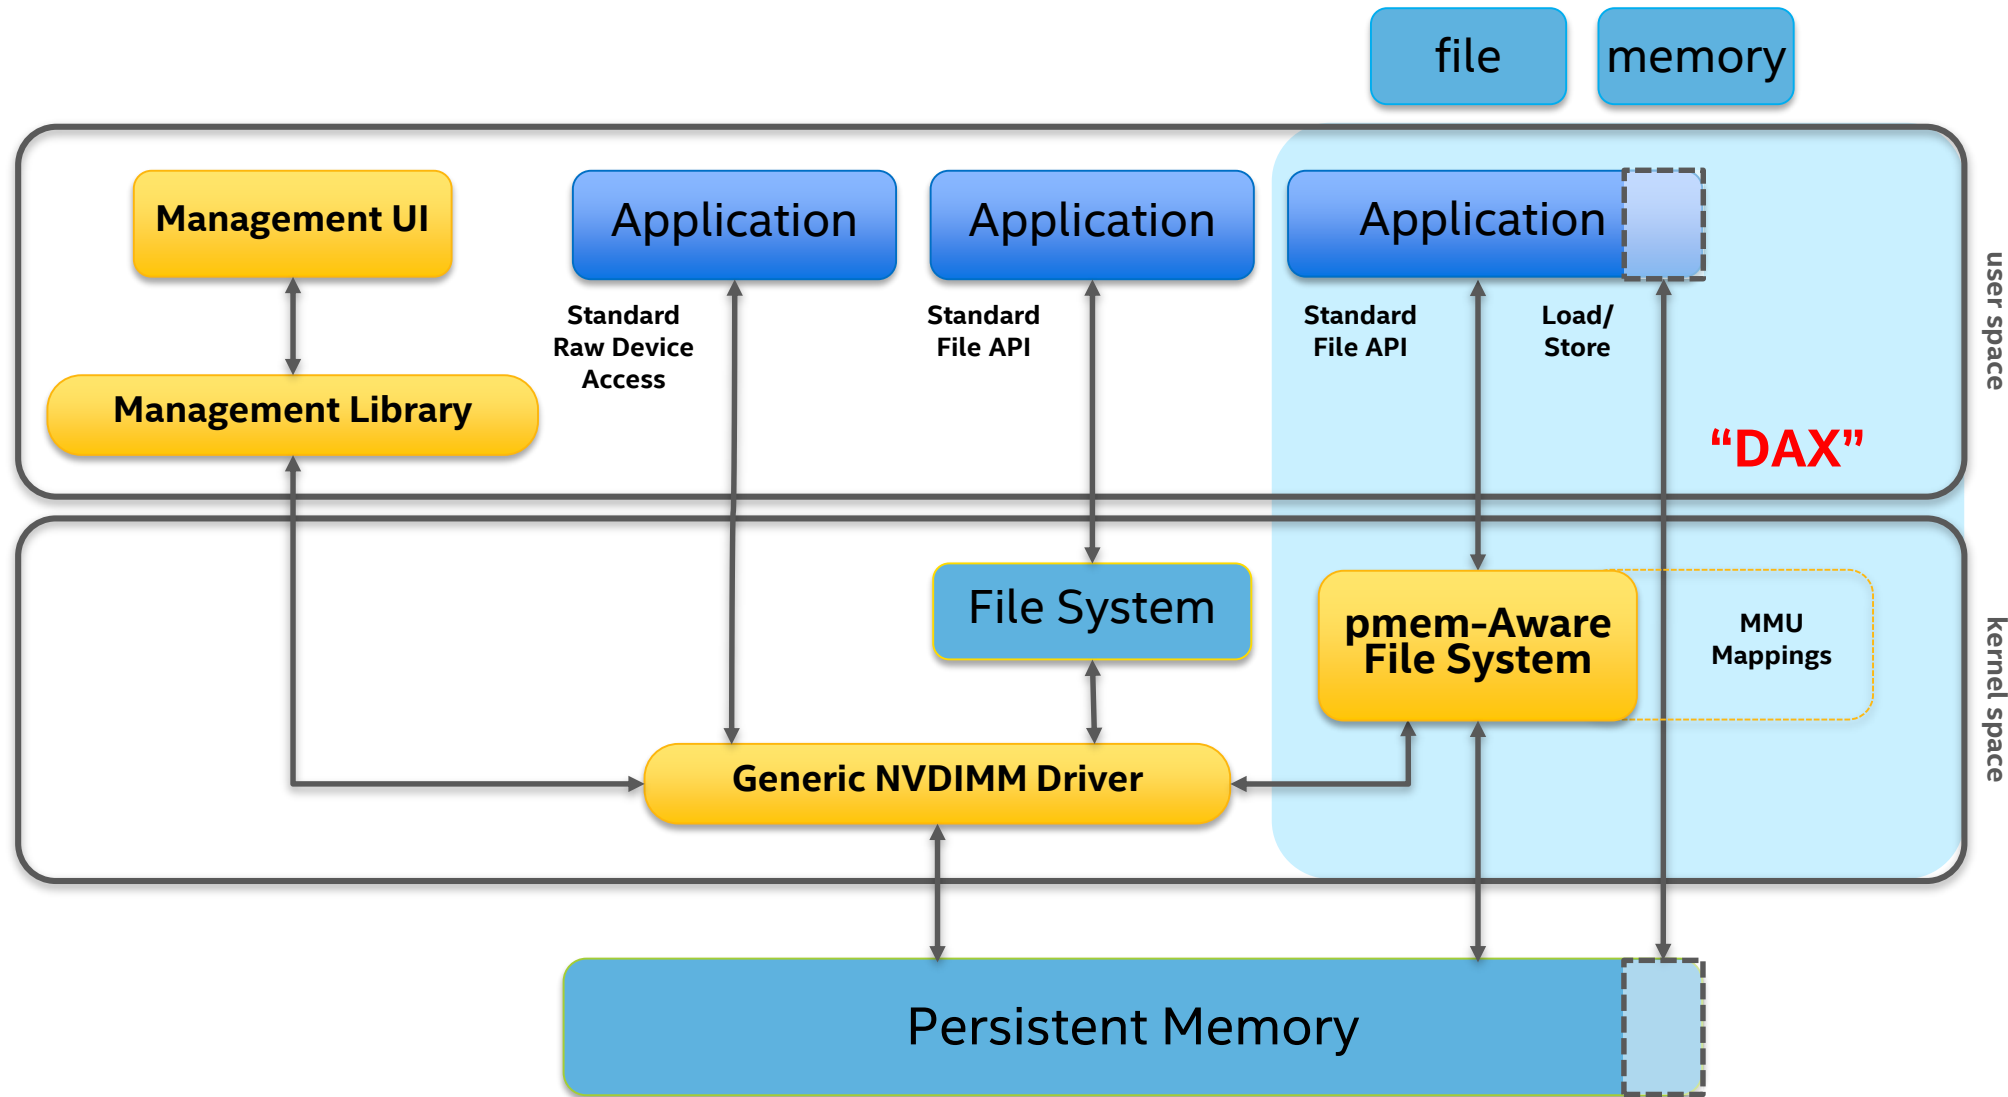
**Piotr Balcer**  
**<piotr.balcer@intel.com>**  
**Intel® Data Center Group**



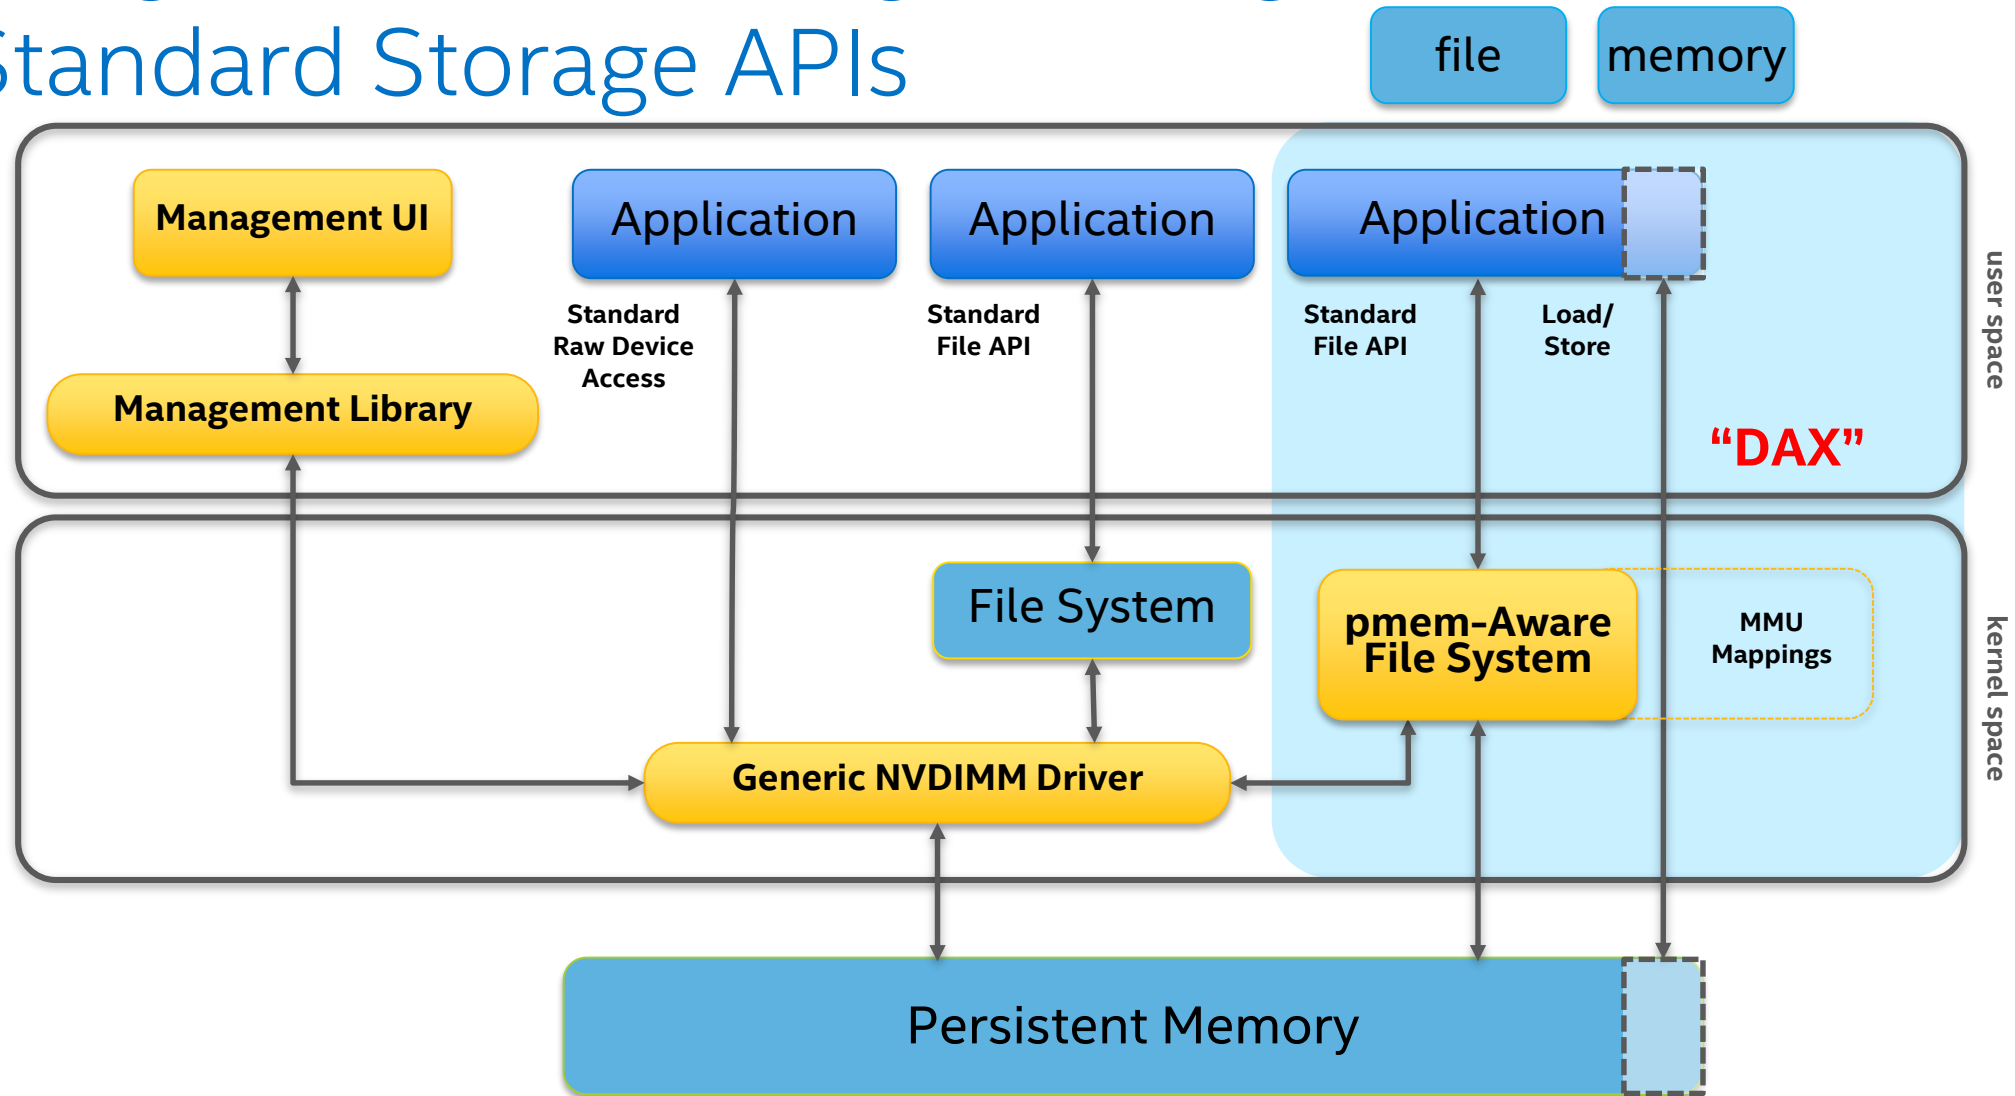
# Agenda

- SNIA NVM Programming Model
  - Block based I/O
  - Memory Mapped I/O
- Understanding power-failure atomicity
- Persistence domain
- Visibility versus Power Fail Atomicity

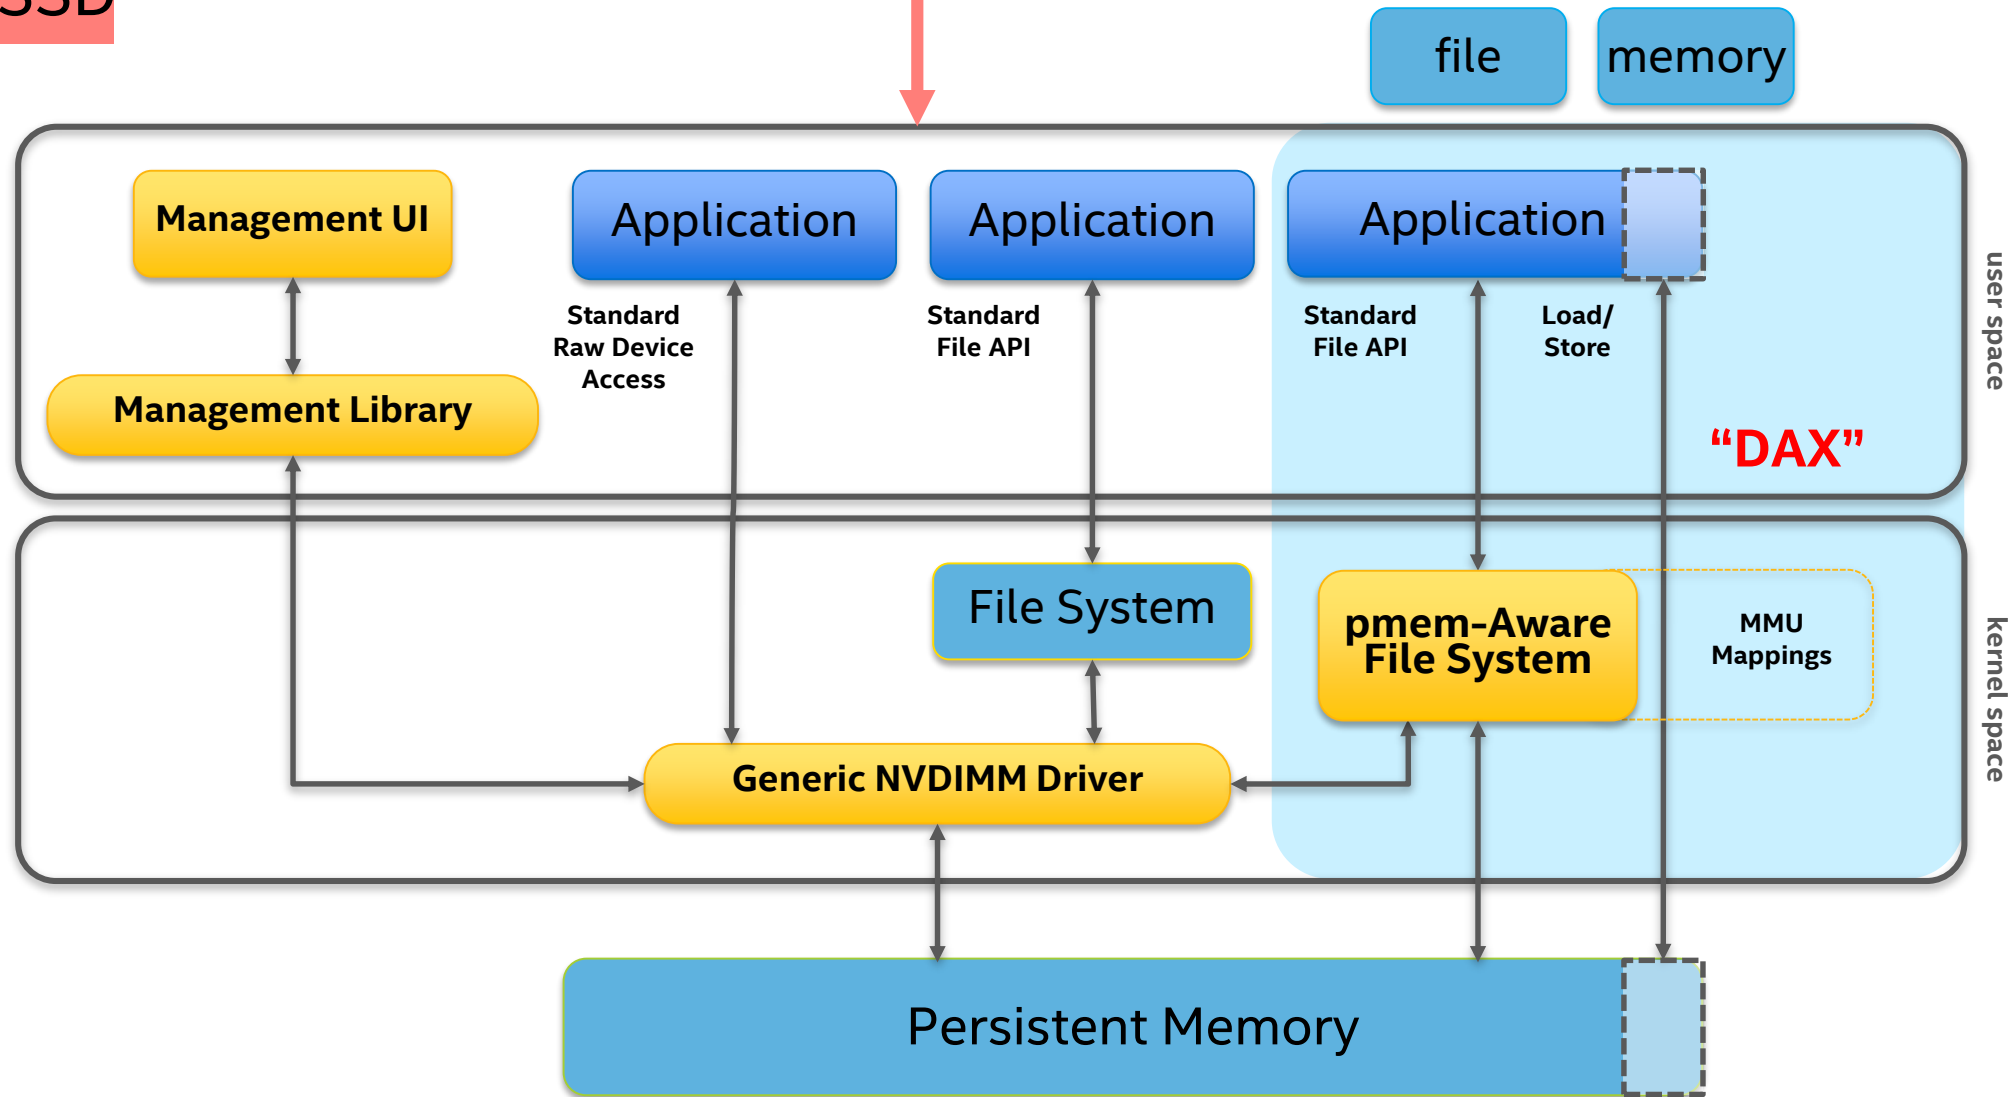
# The SNIA NVM Programming Model



# Don't Forget: The NVM Programming Model Starts With Standard Storage APIs

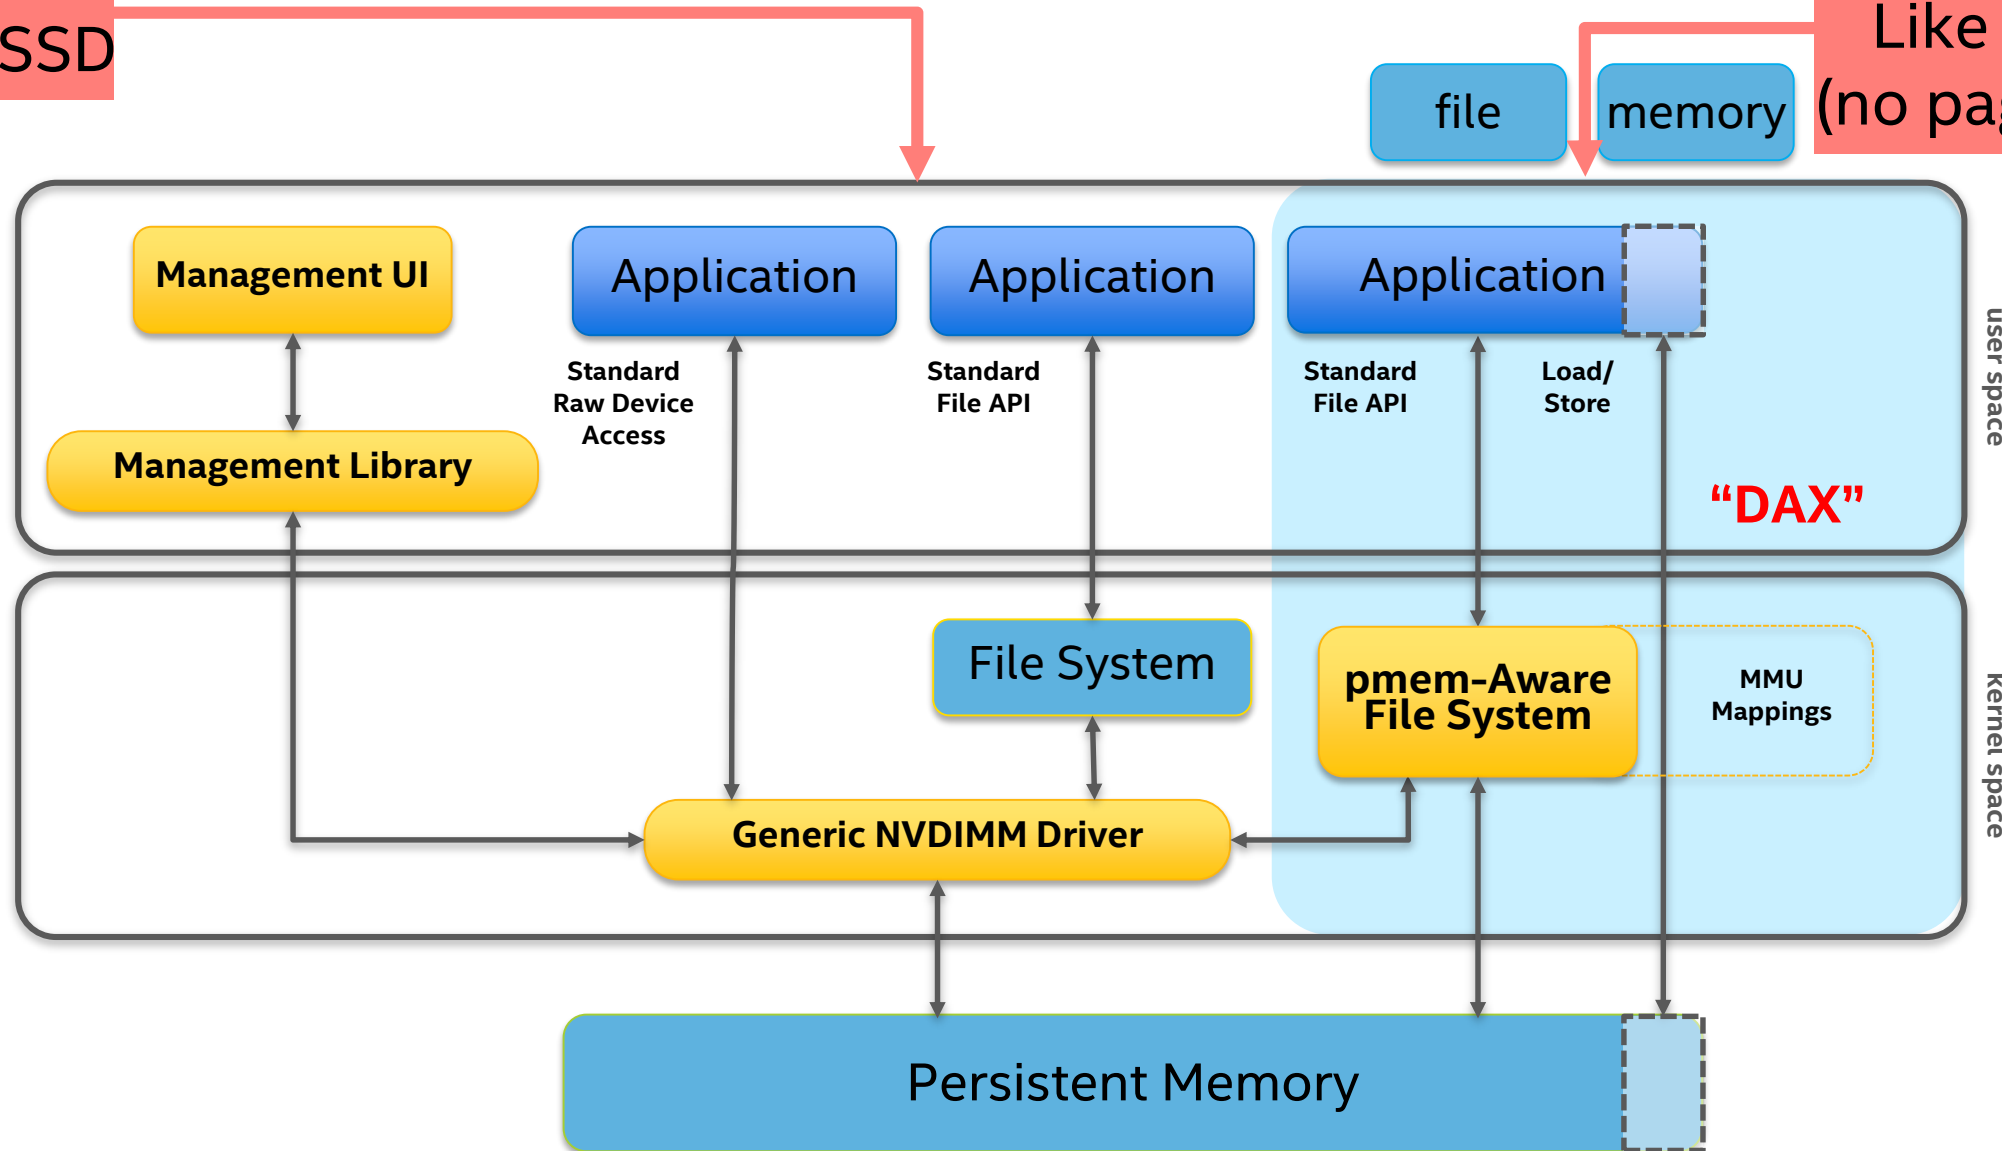


Use PM  
Like an SSD



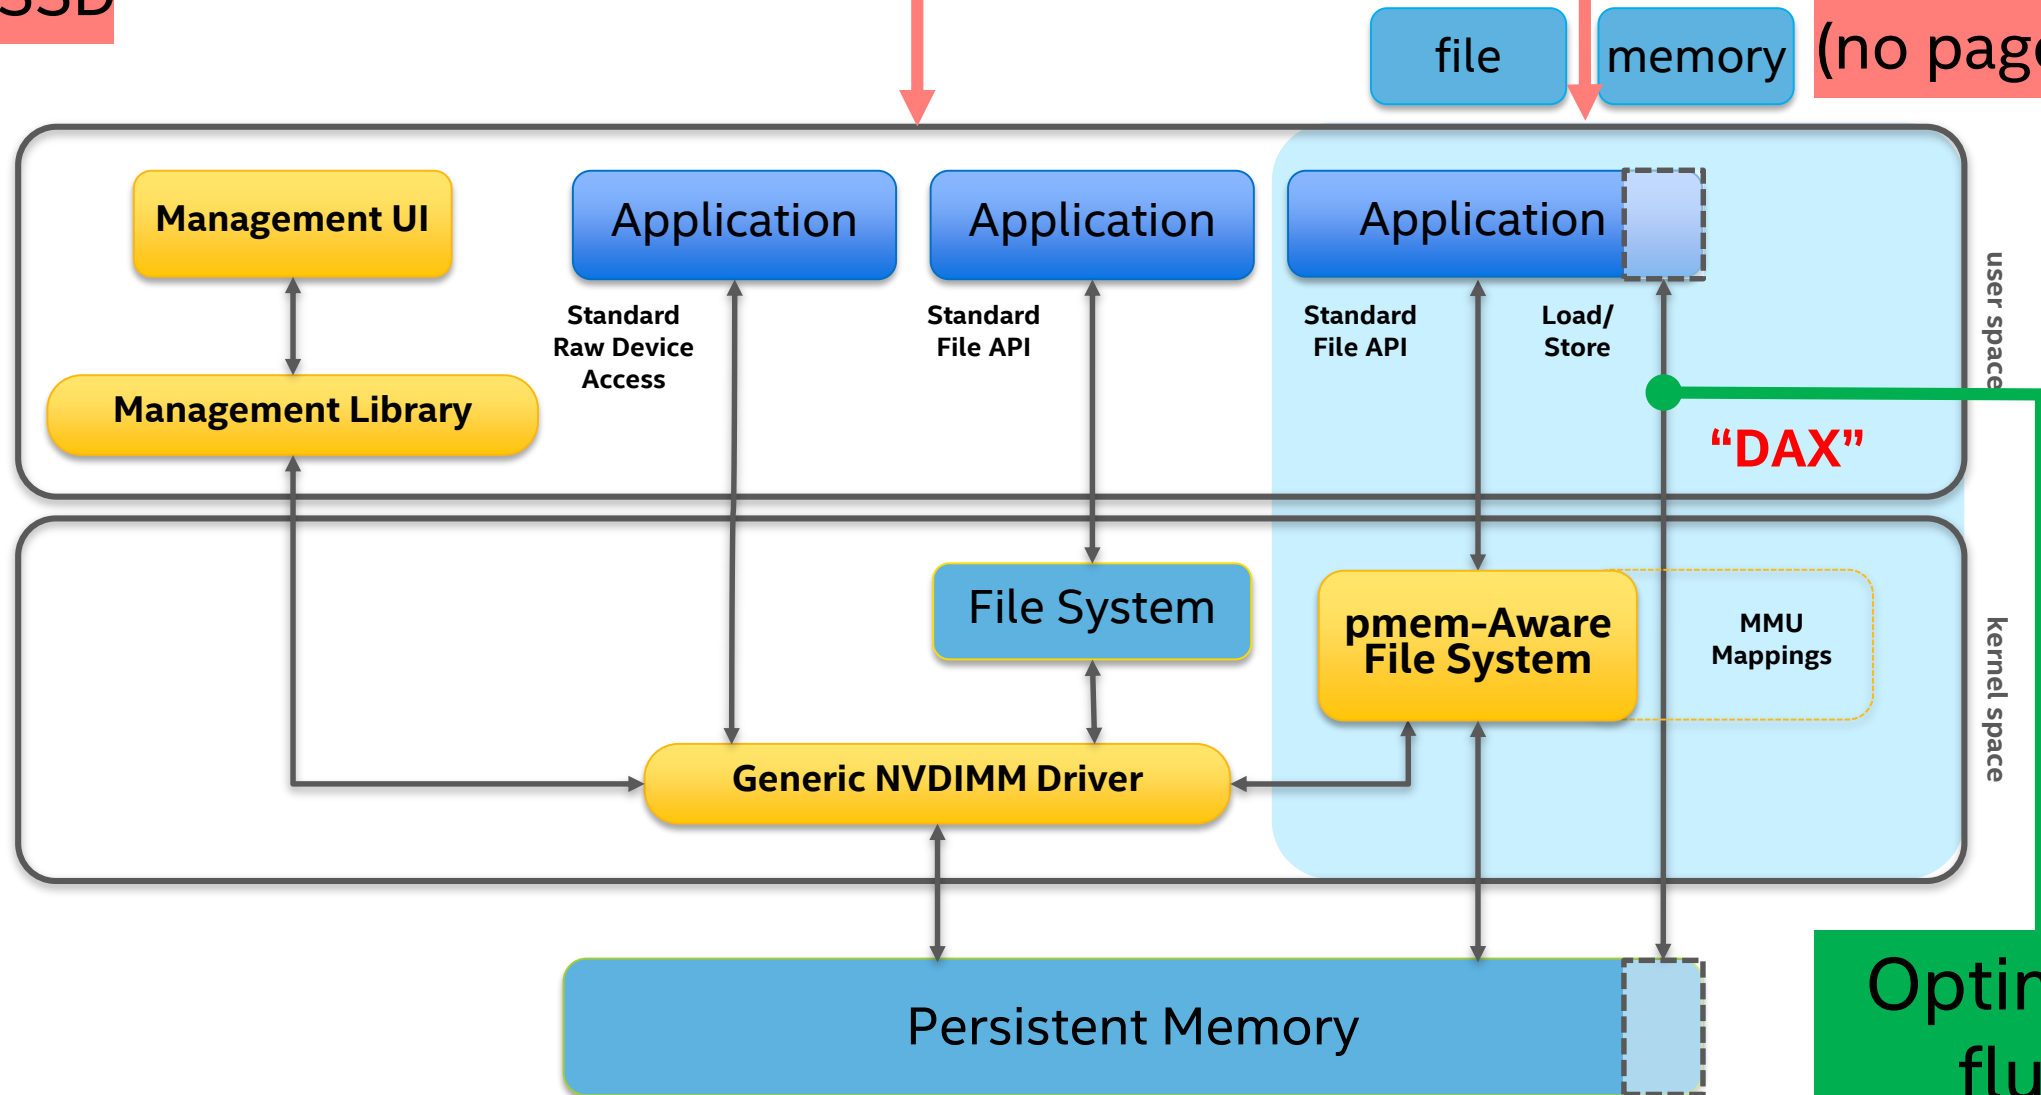
Use PM  
Like an SSD

Use PM  
Like an SSD  
(no page cache)



Use PM  
Like an SSD

Use PM  
Like an SSD  
(no page cache)



# A Programmer's View (mapped files)

```
fd = open("/my/file", O_RDWR);  
...  
base = mmap(NULL, filesize,  
            PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);  
close(fd);  
...  
base[100] = 'X';  
strcpy(base, "hello there");  
*structp = *base_structp;  
...
```

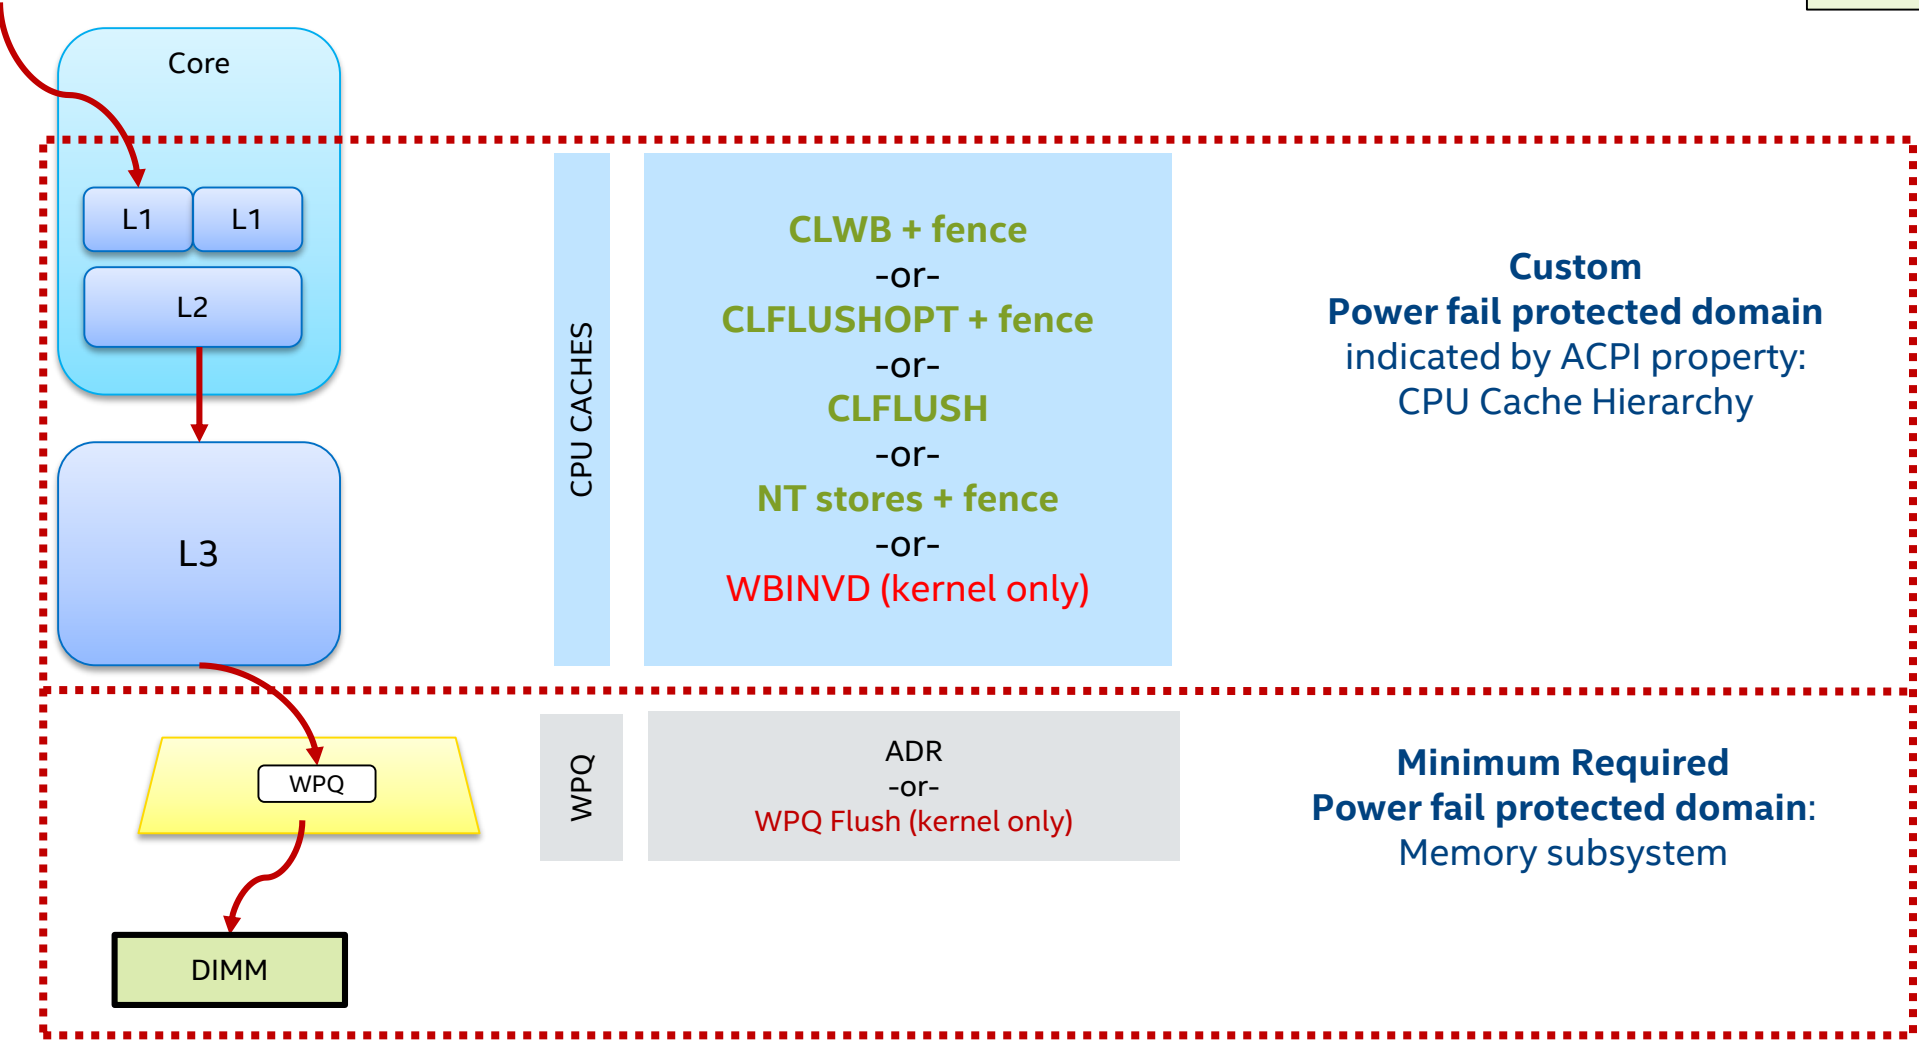
“Load/Store”



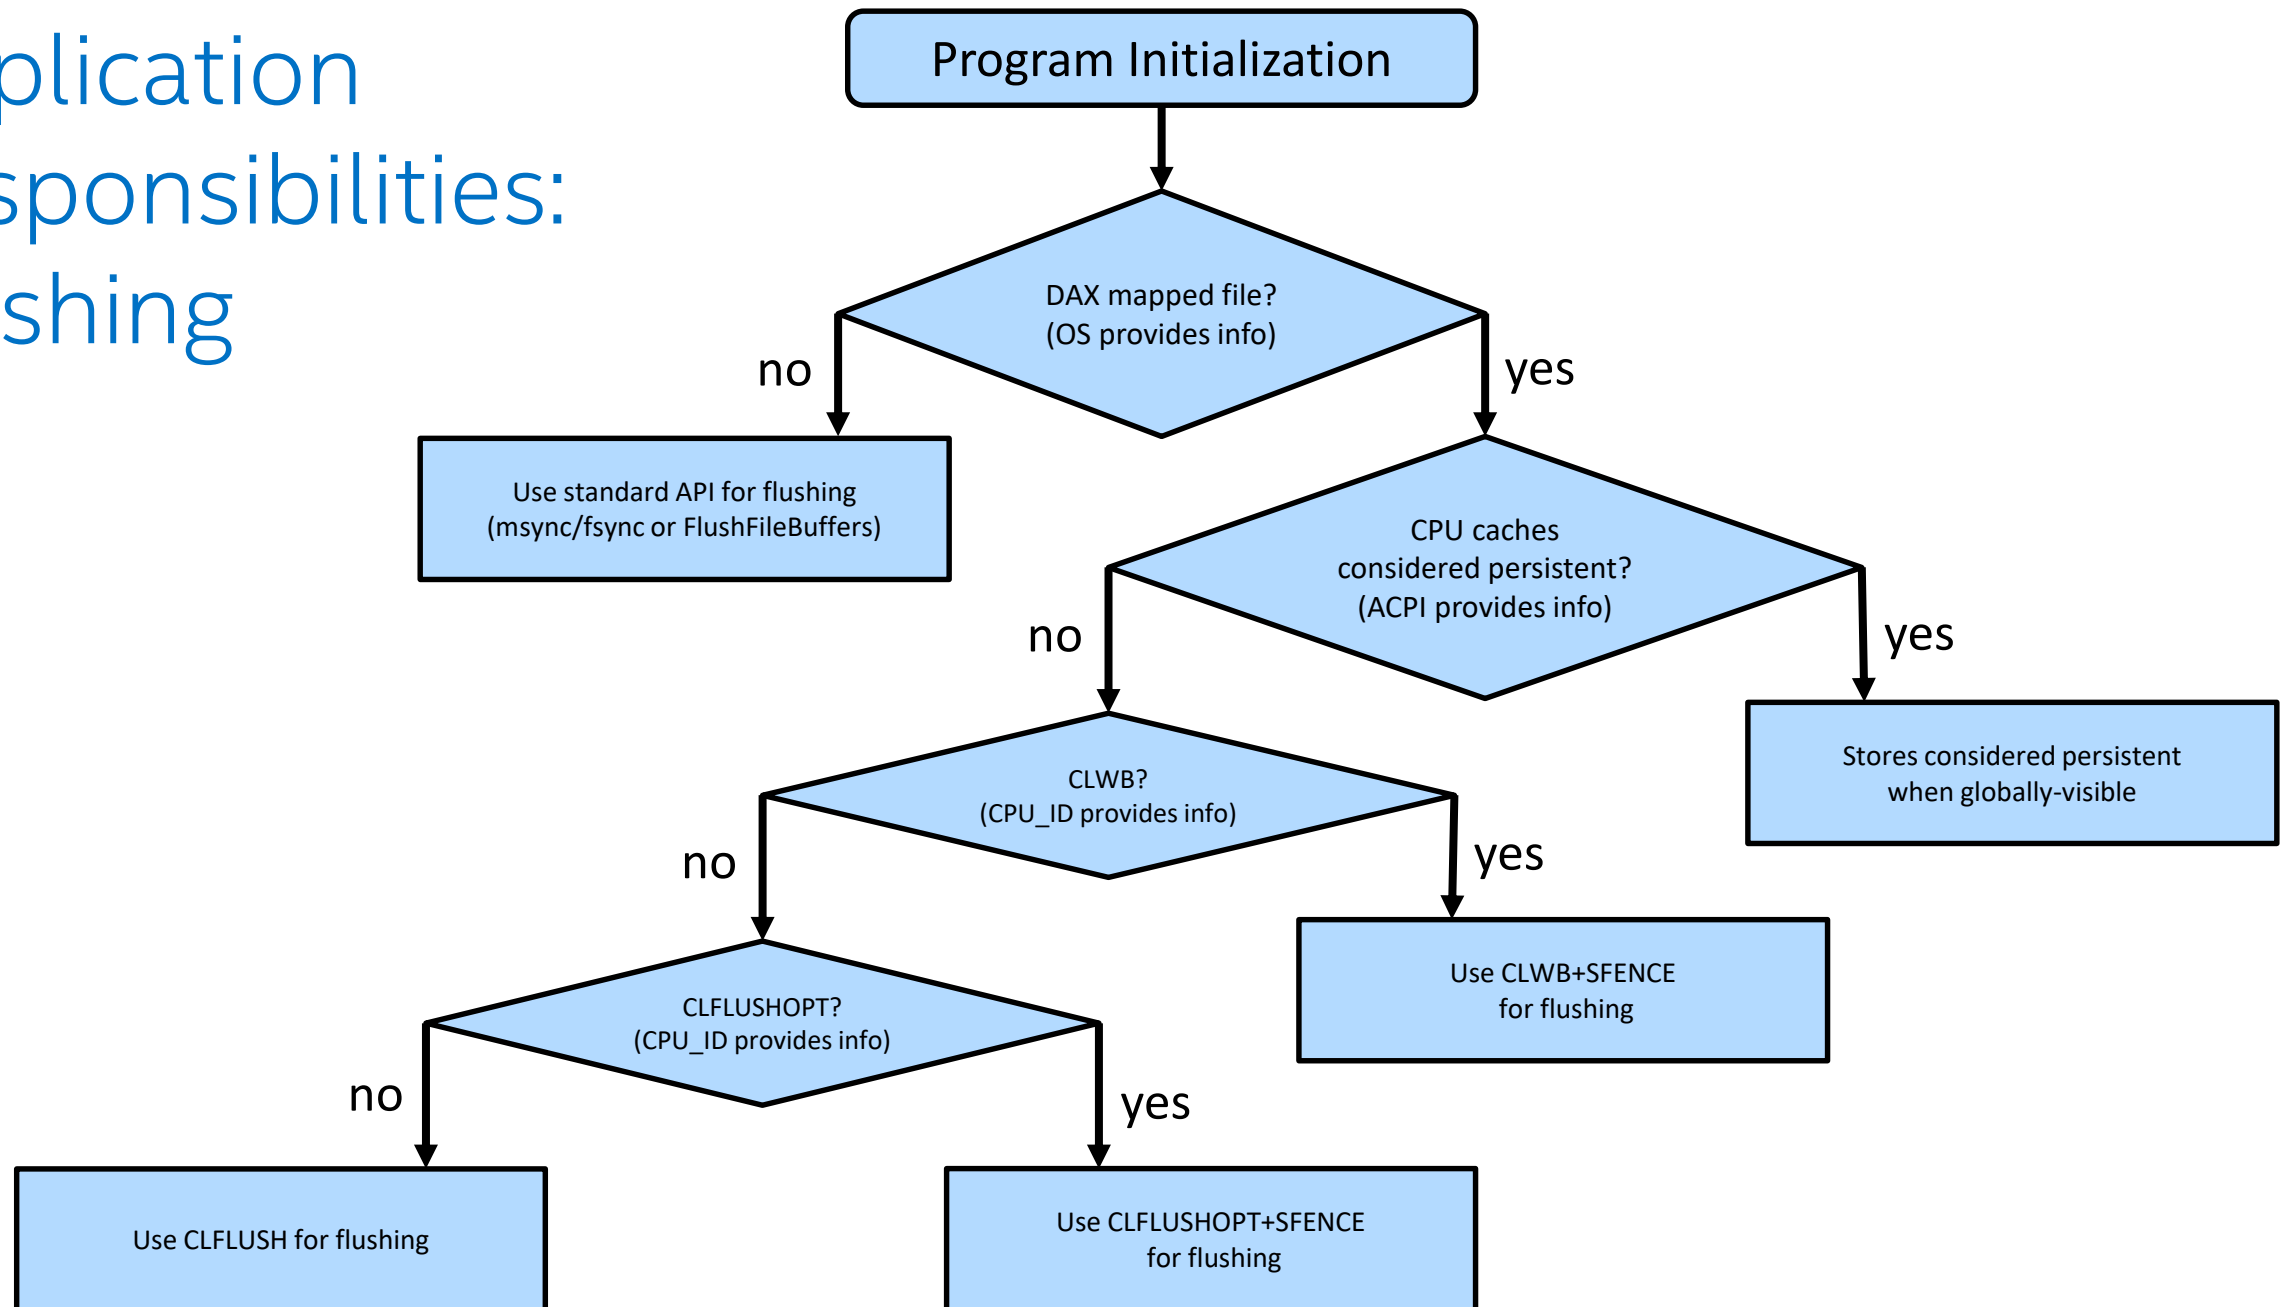
# How the Hardware Works

MOV

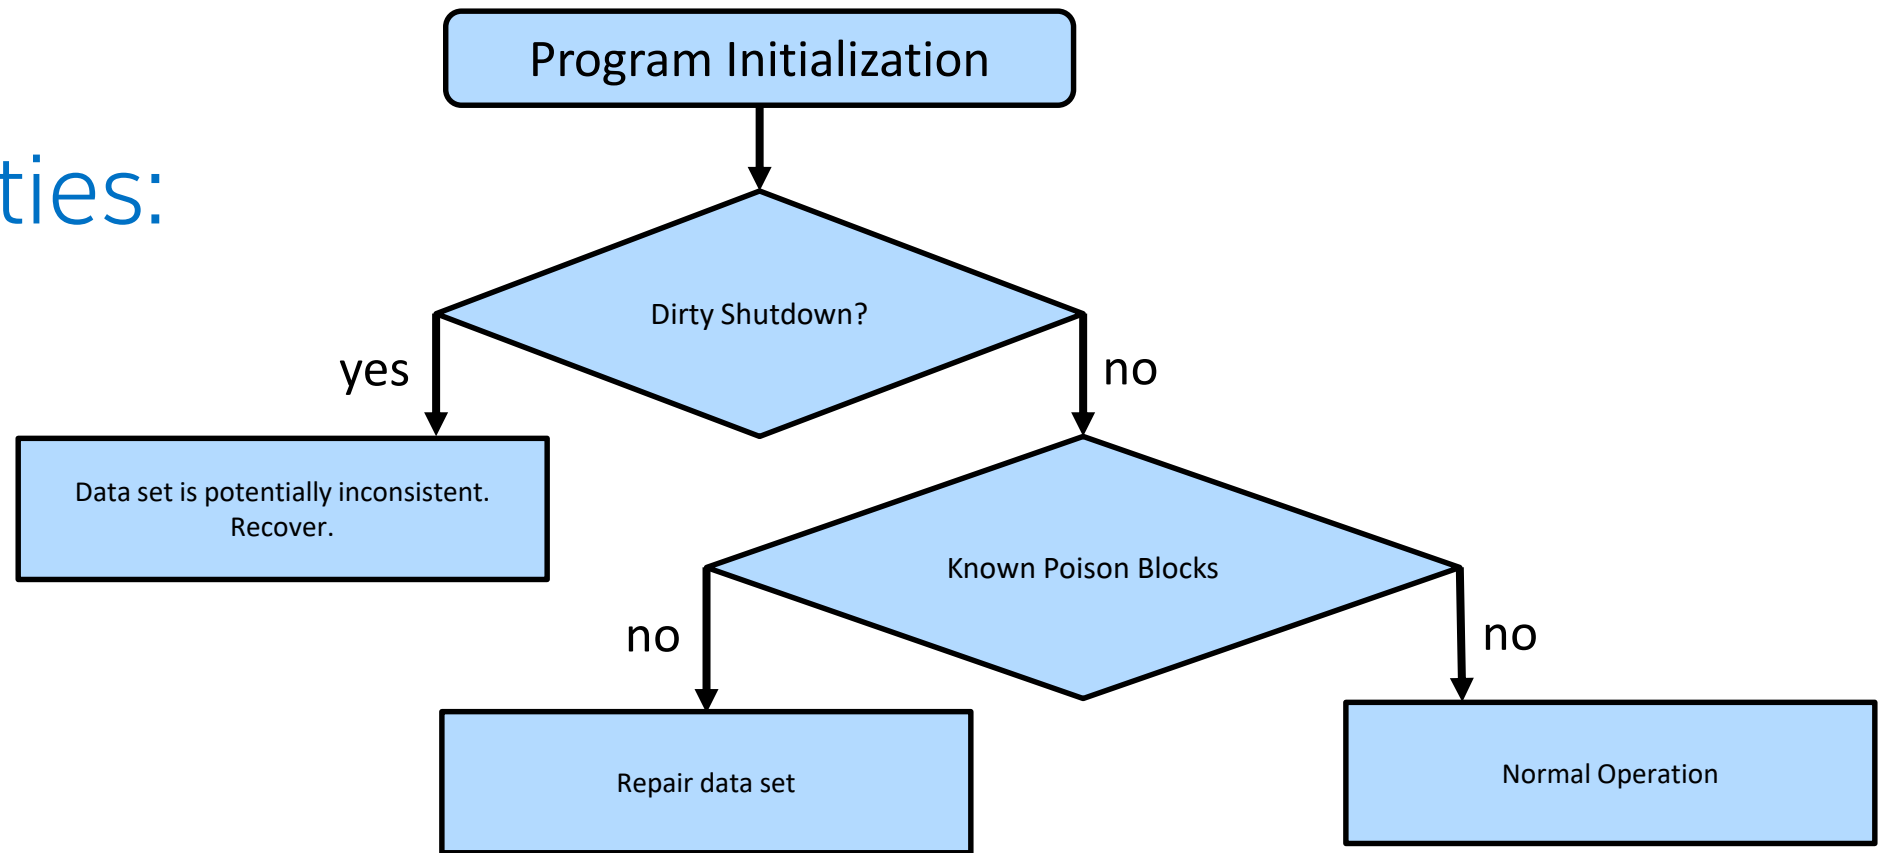
Not shown:  
MCA  
ADR Failure Detection



# Application Responsibilities: Flushing



# Application Responsibilities: Recovery



# Application Responsibilities: Consistency

```
open(...);  
  
mmap(...);  
  
strcpy(pmem, "Hello, World!");  
  
msync(...);
```

← Crash

## Result

1. "\0\0\0\0\0\0\0\0\0\0..."
2. "Hello, w\0\0\0\0\0\0..."
3. "\0\0\0\0\0\0\0\0world!\0"
4. "Hello, \0\0\0\0\0\0\0\0"
5. "Hello, World!\0"



# Application Responsibilities: Consistency

```
open(...);  
  
mmap(...);  
  
strcpy(pmem, "Hello, World!");  
  
pmem_persist(pmem, 14); ← Crash
```

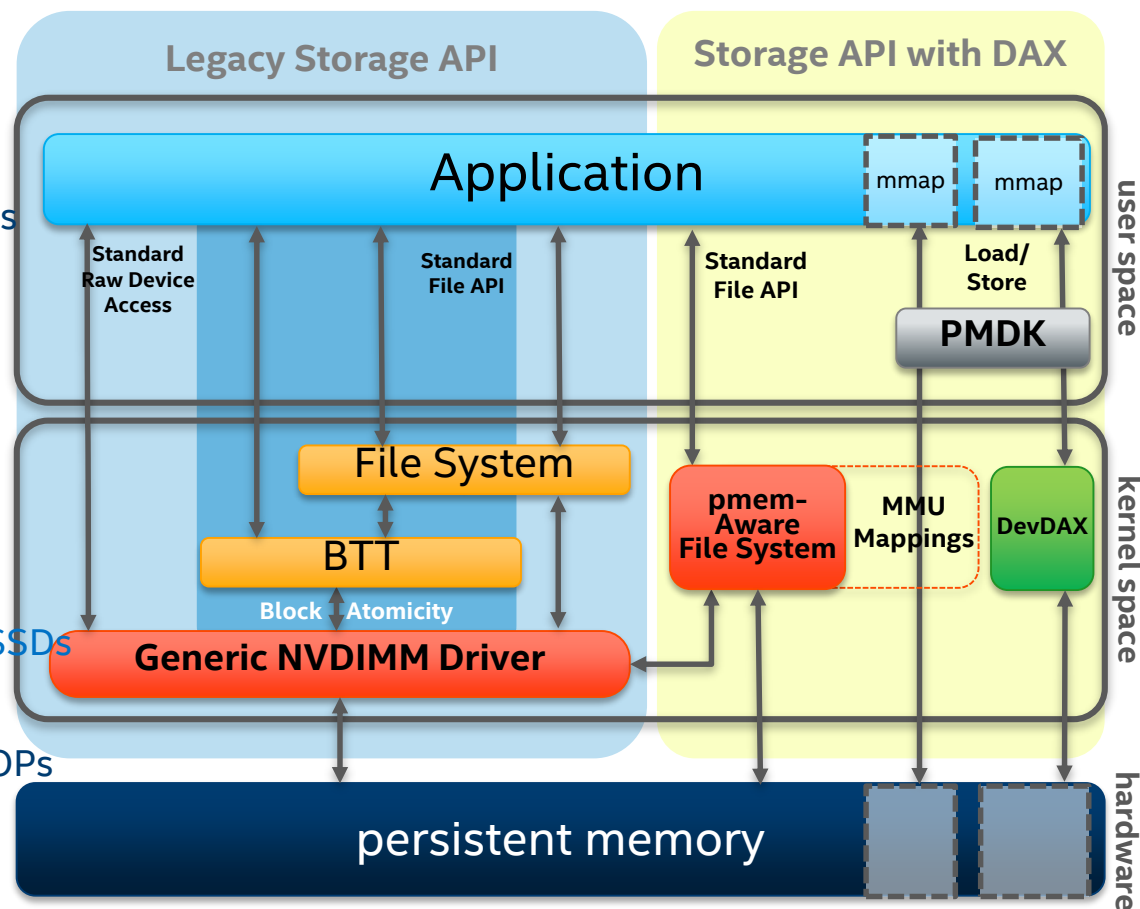
`pmem_persist()` may be faster,  
but is still **not** transactional

## Result

1. "\0\0\0\0\0\0\0\0\0\0..."
2. "Hello, w\0\0\0\0\0\0..."
3. "\0\0\0\0\0\0\0\0world!\0"
4. "Hello, \0\0\0\0\0\0\0\0"
5. "Hello, World!\0"

# Possible ways to access persistent memory

- No Code Changes Required
- Operates in Blocks like SSD/HDD
  - Traditional read/write
  - Works with Existing File Systems
  - Atomicity at block level
  - Block size configurable
    - 4K, 512B\*
- NVDIMM Driver required
  - Support starting Kernel 4.2
- Configured as Boot Device
- Higher Endurance than Enterprise SSDs
- High Performance Block Storage
  - Low Latency, higher BW, High IOPs



- Code changes may be required\*
- Bypasses file system page cache
- Requires DAX enabled file system
  - XFS, EXT4, NTFS
- No Kernel Code or interrupts
- No interrupts
- Fastest IO path possible

\* Code changes required for load/store direct access if the application does not already support this.

\*Requires Linux

# Visibility versus Power Fail Atomicity

Feature	Atomicity
Atomic Store	8 byte powerfail atomicity Much larger visibility atomicity
TSX	Programmer must comprehend XABORT, cache flush can abort
LOCK CMPXCHG	Non-blocking algorithms depend on CAS, but CAS doesn't include flush to persistence

Software must implement all atomicity beyond 8 bytes for pmem  
Transactions are fully up to software

# If caches are not flush on failure...

- Can't easily use `compare_and_swap` / `fetch_and_add` on Persistent Memory resident variables
- Can't use Hardware Transactional Memory (TSX) on Persistent Memory
- Must manually flush all data after writing

# If caches are flush on failure...

- No need to flush data
- But applications still need to do their own transactions
  - Can use HTM/TSX for that, must include a software fallback in case hardware transaction fails



# PMEM reference counter – BAD example

```
struct my_object {  
    uint64_t refcount;  
    type some_resource;  
};
```

No decision based on this value in this thread...

```
static void object_ref(struct my_object *object) { /* refcount visible = 0    persistent = 0 */  
    __sync_fetch_and_add(&object->refcount, 1); /* visible = 1    persistent = ? */  
    persist(&object->refcount, sizeof(object->refcount)); /* visible = 1    persistent = 1 */  
}
```

Decision is made based on visible but not persistent value

```
static void object_deref(struct my_object *object) { /* visible = 1    persistent = 1 */  
    if (__sync_sub_and_fetch(&object->refcount, 1) == 0) { /* visible = 0    persistent = ? */  
        delete_some_resource(object->some_resource); /* visible = 0    persistent = ? */  
    }  
    persist(&object->refcount, sizeof(object->refcount)); /* visible = 0    persistent = 0 */  
}
```

# PMEM reference counter – GOOD example

```
struct my_object {  
    uint64_t refcount;  
    type some_resource;  
};
```

No decision based on this value in this thread...

```
static void object_ref(struct my_object *object) { /* refcount visible = 0      persistent = 0 */  
    __sync_fetch_and_add(&object->refcount, 1); /*      visible = 1      persistent = ? */  
    persist(&object->refcount, sizeof(object->refcount)); /* visible = 1      persistent = 1 */  
}
```

Decision is based on a known persistent value

```
static void object_deref(struct my_object *object) { /*      visible = 1      persistent = 1 */  
    if (__sync_sub_and_fetch(&object->refcount, 1) == 0) { /*      visible = 0      persistent = ? */  
        persist(&object->refcount, sizeof(object->refcount)); /* visible = 0      persistent = 0 */  
        delete_some_resource(object->some_resource); /*      visible = 0      persistent = 0 */  
    }  
}
```

Atomic variables need to be read and flushed before making any decisions/calculations with them to ensure that the action is taken on a value that is known to have been persistent at some point.





谢谢



